

Olivia Hsu

owhsu@stanford.edu | [Personal Website](#)

 [LinkedIn](#) |  [weiya711](#) |  [Google Scholar](#)

EDUCATION

Stanford University

Ph.D. in Computer Science (CS)

Jan 2026
Stanford, USA

- Advisors: Kunle Olukotun and Fredrik Kjolstad
- Thesis: [Programming Systems for Sparse Accelerators](#)

Stanford University

M.S. in Computer Science (CS)

Sep 2025
Stanford, USA

University of California, Berkeley

B.S. in Electrical Engineering and Computer Science (EECS)

May 2019
Berkeley, USA

ACADEMIC AND RESEARCH EXPERIENCE

• Carnegie Mellon University

Assistant Professor

May 2026 – Present
Pittsburgh, USA

- Assistant professor in Electrical and Computer Engineering and, by courtesy, Computer Science.
- Member of the CMU Parallel Data Lab (PDL)  and Catalyst Lab .

• Carnegie Mellon University

Courtesy Adjunct Assistant Professor

June 2025 – Apr 2026
Pittsburgh, USA

• Stanford Pervasive Parallelism Lab

Postdoc and PhD Student Researcher

Jan 2020 – Apr 2026
Stanford, USA

- Researching the automatic compilation and mapping of sparse applications to domain-specific architectures, reconfigurable dataflow architectures, and hardware accelerators under professors Kunle Olukotun and Fredrik Kjolstad.
- During my postdoc, I advised research projects [1, 2]
- During my PhD, I led research projects [7, 13, 16, 26, 10, 8], mentored research projects [12, 9, 25, 5], and have helped on research projects [11, 0, 3].

• UC Berkeley Group Matching Project

Student Collaborator

Aug 2020 – Aug 2021
Remote

- Helped develop a scalable and inclusive study group matching process for student collaboration under Professor Gireeja Ranade, which led to [14].

• Stanford VLSI Group

PhD Student Researcher

Sep 2019 – Jan 2020
Stanford, USA

- Researching fully digital ADC-based high-speed link equalization techniques under Professor Mark Horowitz.

• Berkeley Wireless Research Center

Undergraduate Researcher

June 2017 – Aug 2019
Berkeley, USA

- Designed photonic devices using the Berkeley Photonics Generator (BPG) and digital circuits for high-speed communication in Python and Verilog for an ultrasound receive-array system fabricated in the Global Foundries 45nm RFSOI process.
- Built the lab setup and tested the sensitivity of silicon photonic ring modulators when disturbed with ultrasonic waves.

Refereed Conference and Journal Papers

- [1] Rubens Lacouture, Nathan Zhang, Ritvik Sharma, Marco Siracusa, Fredrik Kjolstad, Kunle Olukotun, and **Olivia Hsu**. “FuseFlow: A Fusion-Centric Compilation Framework for Sparse Deep Learning on Streaming Dataflow”. In: *Proceedings of the 31st ACM International Conference on Architectural Support for Programming Languages and Operating Systems*. 2026. URL: [🔗](#)
- [2] Gina Sohn, Genghan Zhang, Konstantin Hossfeld, Jungwoo Kim, Nathan Sobotka, Nathan Zhang, **Olivia Hsu**, and Kunle Olukotun. “Streaming Tensor Programs: A Streaming Abstraction for Dynamic Parallelism”. In: *Proceedings of the 31st ACM International Conference on Architectural Support for Programming Languages and Operating Systems*. 2026. URL: [🔗](#)
- [3] Marco Siracusa, **Olivia Hsu**, Victor Soria-Pardos, Joshua Randall, Arnaud Grasset, Eric Biscondi, Doug Joseph, Randy Allen, Fredrik Kjolstad, Miquel Moretó Planas, and Adrià Armejach. “Ember: A Compiler for Efficient Embedding Operations on Decoupled Access-Execute Architectures”. In: *Proceedings of the 2026 IEEE/ACM International Symposium on Code Generation and Optimization (CGO)*. Sydney, Australia, 2026.
- [4] Kalhan Koul*, **Olivia Hsu***, Yuchen Mei, Sai Gautham Ravipati, Maxwell Strange, Jackson Melchert, Alex Carsello, Taeyoung Kong, Po-Han Chen, Huifeng Ke, Keyi Zhang, Qiaoyi Liu, Gedeon Nyengele, Zhouhua Xie, Akhilesh Balasingam, Jayashree Adivarahan, Ritvik Sharma, Christopher Torng, Joel S. Emer, Fredrik Kjolstad, Mark Horowitz, and Priyanka Raina. “Onyx: A 12-nm Programmable Accelerator for Dense and Sparse Applications”. In: *IEEE Journal of Solid-State Circuits* (2025).
*Note: * signifies equal contribution.*
- [5] Genghan Zhang, Weixin Liang, **Olivia Hsu**, and Kunle Olukotun. “Adaptive Self-Improvement LLM Agentic System for ML Library Development”. In: *International Conference on Machine Learning (ICML)*. 2025.
- [6] Kalhan Koul*, Zhouhua Xie*, Maxwell Strange, Sai Gautham Ravipati, Bo Wun Cheng, **Olivia Hsu**, Po-Han Chen, Mark Horowitz, Fredrik Kjolstad, and Priyanka Raina. “Designing Programmable Accelerators for Sparse Tensor Algebra 2025”. In: *IEEE Micro*. Hot Chips Special Issue. 2025. *Note: * signifies equal contribution.*
- [7] **Olivia Hsu**, Alexander Rucker, Tian Zhao, Varun Desai, Kunle Olukotun, and Fredrik Kjolstad. “Stardust: Compiling Sparse Tensor Algebra to a Reconfigurable Dataflow Architecture”. In: *Proceedings of the 23rd ACM/IEEE International Symposium on Code Generation and Optimization*. CGO ’25. Las Vegas, NV, USA: Association for Computing Machinery, 2025. URL: [🔗](#)
- [8] Kalhan Koul, Maxwell Strange, Jackson Melchert, Alex Carsello, Yuchen Mei, **Olivia Hsu**, Taeyoung Kong, Po-Han Chen, Huifeng Ke, Keyi Zhang, Qiaoyi Liu, Gedeon Nyengele, Akhilesh Balasingam, Jayashree Adivarahan, Ritvik Sharma, Zhouhua Xie, Christopher Torng, Joel Emer, Fredrik Kjolstad, Mark Horowitz, and Priyanka Raina. “Onyx: A Programmable Accelerator for Sparse Tensor Algebra”. In: *2024 IEEE Hot Chips 36 Symposium (HCS)*. Los Alamitos, CA, USA: IEEE Computer Society, Aug. 2024. URL: [🔗](#)
- [9] Genghan Zhang, **Olivia Hsu**, and Fredrik Kjolstad. “Compilation of Modular and General Sparse Workspaces”. In: *Proc. ACM Program. Lang.* 8.PLDI (June 2024). URL: [🔗](#)
- [10] Kalhan Koul, Maxwell Strange, Jackson Melchert, Alex Carsello, Yuchen Mei, **Olivia Hsu**, Taeyoung Kong, Po-Han Chen, Huifeng Ke, Keyi Zhang, Qiaoyi Liu, Gedeon Nyengele, Akhilesh Balasingam, Jayashree Adivarahan, Ritvik Sharma, Zhouhua Xie, Christopher Torng, Joel Emer, Fredrik Kjolstad, Mark Horowitz, and Priyanka Raina. “Onyx: A 12nm 756 GOPS/W Coarse-Grained Reconfigurable Array for Accelerating Dense and Sparse Applications”. In: *2024 IEEE Symposium on VLSI Technology and Circuits (VLSI Technology and Circuits)*. 2024. URL: [🔗](#)
- [11] Erik Orm Hellsten, Artur Souza, Johannes Lenfers, Rubens Lacouture, **Olivia Hsu**, Adel Ejje, Fredrik Kjolstad, Michel Steuwer, Kunle Olukotun, and Luigi Nardi. “BaCO: A Fast and Portable Bayesian Compiler Optimization Framework”. In: *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 4*. ASPLOS ’23. Vancouver, BC, Canada: Association for Computing Machinery, 2024. URL: [🔗](#)
- [12] Manya Bansal, **Olivia Hsu**, Kunle Olukotun, and Fredrik Kjolstad. “Mosaic: An Interoperable Compiler for Tensor Algebra”. In: *Symposium on Programming Languages Design and Implementation (PLDI)*. **Distinguished Paper**. 2023. URL: [🔗](#)

- [13] **Olivia Hsu**, Maxwell Strange, Ritvik Sharma, Jaeyeon Won, Kunle Olukotun, Joel S. Emer, Mark A. Horowitz, and Fredrik Kjolstad. “The Sparse Abstract Machine”. In: *International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*. ASPLOS 2023. Vancouver, BC, Canada: Association for Computing Machinery, 2023. URL: [🔗](#)
- [14] Sumer Kohli, Neelesh Ramachandran, Ana Tudor, Gloria Tumushabe, **Olivia Hsu**, and Gireeja Ranade. “Inclusive Study Group Formation at Scale”. In: *Proceedings of the 54th ACM Technical Symposium on Computer Science Education V. 1*. SIGCSE 2023. Toronto ON, Canada: Association for Computing Machinery, 2023. URL: [🔗](#)
- [15] Panagiotis Zarkos, Sidney Buchbinder, Christos Adamopoulos, Sarika Madhvapathy, **Olivia Hsu**, Jake Whinnery, Pavan Bhargava, and Vladimir Stojanovic. “Fully Integrated Electronic-Photonic Ultrasound Receiver Array for Endoscopic Applications in a Zero-Change 45-nm CMOS-SOI Process”. In: *IEEE Journal of Solid-State Circuits* 58.6 (2023). URL: [🔗](#)
- [16] Rawn Henry*, **Olivia Hsu***, Rohan Yadav, Stephen Chou, Kunle Olukotun, Saman Amarasinghe, and Fredrik Kjolstad. “Compilation of sparse array programming models”. In: *Proceedings of the ACM on Programming Languages (OOPSLA)*. 2021. URL: [🔗](#) Note: * signifies equal contribution.
- [17] Panagiotis Zarkos, Sidney Buchbinder, Christos Adamopoulos, Sarika Madhvapathy, **Olivia Hsu**, Jake Whinnery, Pavan Bhargava, and Vladimir Stojanovic. “Fully Integrated Electronic-Photonic Ultrasound Receiver Array for Endoscopic Imaging Applications in a Zero-Change 45nm CMOS-SOI Process”. In: *2021 Symposium on VLSI Circuits*. 2021. URL: [🔗](#)
- [18] Panagiotis Zarkos, Sidney Buchbinder, Christos Adamopoulos, **Olivia Hsu**, Sarika Madhvapathy, Jake Whinnery, Pavan Bhargava, and Vladimir Stojanovic. “Monolithically Integrated Electronic-Photonic Ultrasound Receiver Using Microring Resonator”. In: *2021 Conference on Lasers and Electro-Optics (CLEO)*. 2021.
- [19] Panagiotis Zarkos, **Olivia Hsu**, and Vladimir Stojanovic. “Ring Resonator Based Ultrasound Detection in a Zero-Change Advanced CMOS-SOI Process”. In: *2019 Conference on Lasers and Electro-Optics (CLEO)*. 2019. URL: [🔗](#)

Refereed Workshop Papers

- [20] Genghan Zhang, Weixin Liang, **Olivia Hsu**, and Kunle Olukotun. *Adaptive Self-improvement LLM Agentic Aystem for ML Library Development*. Workshop on Efficient Systems for Foundation Models (ES-FoMo) at ICML 2025. 2025.
- [21] Genghan Zhang, Weixin Liang, **Olivia Hsu**, and Kunle Olukotun. *Adaptive Self-improvement LLM Agentic Aystem for ML Library Development*. MLSys Young Professionals Workshop co-located with MLSys 2025. 2025.
- [22] Genghan Zhang, Weixin Liang, **Olivia Hsu**, and Kunle Olukotun. *Adaptive Self-improvement LLM Agentic Aystem for ML Library Development*. Workshop on Deep Learning for Code (DL4C) @ ICLR 2025. **Best Paper**. 2025.
- [23] Genghan Zhang, Weixin Liang, **Olivia Hsu**, and Kunle Olukotun. *Adaptive Self-improvement LLM Agentic Aystem for ML Library Development*. Workshop on Reasoning and Planning for LLMs @ ICLR 2025. 2025.
- [24] Marco Siracusa, **Olivia Hsu**, Victor Soria-Pardos, Joshua Randall, Arnaud Grasset, Eric Biscondi, Doug Joseph, Randy Allen, Fredrik Kjolstad, Miquel Moretó Planas, and Adrià Armejach. *An Architecture and a Compiler to Accelerate Embedding Operations in Machine Learning Models*. Compilers for Machine Learning (C4ML) Workshop co-located with CGO 2025. 2025.
- [25] Rubens Lacouture, **Olivia Hsu**, Kunle Olukotun, and Fredrik Kjolstad. *Challenges with Hardware-Software Co-design for Sparse Machine Learning on Streaming Dataflow*. Workshop on Programming Languages and Architecture (PLARCH) co-located with FCRC/ISCA/PLDI. 2023.
- [26] **Olivia Hsu**, Maxwell Strange, Kunle Olukotun, Mark Horowitz, and Fredrik Kjolstad. *Designing a Dataflow Hardware Accelerator with an Abstract Machine*. Workshop on Languages, Tools, and Techniques for Accelerator Design (LATTE) co-located with ASPLOS. 2023.

TEACHING EXPERIENCE

- **Stanford Agile Hardware Design (AHA) Tutorial at MICRO 2024**
Presenter

Nov 2024
Austin, USA

- Teaching the sparse application flow for the Stanford AHA project, which includes the front-end sparse Custard compiler [13, 4], the sparse abstract machine (SAM) dataflow intermediate representation [13], and sparse optimizations for the Onyx coarse-grained reconfigurable array (CGRA) [8, 10, 4].
- **Stanford University CS 143: Compilers** Spring 2023, Spring 2024
Course Assistant Stanford, USA
 - Taught under professor Fredrik Kjolstad.
 - Duties include running office hours, managing course logistics, developing written assignments, debugging exams, writing exam problems, and homework and exam grading.
- **Stanford University CS149: Parallel Programming** Fall 2021
Course Assistant Stanford, USA
 - Taught under professors Kayvon Fatahalian and Kunle Olukotun.
 - Duties included running office hours, debugging exams, and exam and homework grading.
- **UC Berkeley Decal EE 198: Hands-on PCB Engineering (HOPE)** Fall 2018
Student Instructor Berkeley, USA
 - Co-created and developed the curriculum for the first-ever offering of this course, an undergraduate-led class on printed circuit board design affiliated with the UC Berkeley IEEE student chapter.
 - Taught the course to approximately 30 students, teaching these students how to read design documentation, design circuits, use PCB design tools, package design files for PCB fabrication, and populate and test their own PCB designs.
- **UC Berkeley EECS 16A: Designing Information Devices and Systems I** Fall 2016 – Fall 2017
Undergraduate Student Instructor (UGSI) Berkeley, USA
 - Administrative and Lab UGSI (Fall 2016) and Head and Lab UGSI (Spring – Fall 2017) for the introductory electrical engineering course. The largest offering I taught had 850+ students, and I was in charge of 30+ teaching assistants.
 - Duties included managing course administration and policies, responding to all student emails related to the course, teaching and debugging labs, creating exam problems, running office hours, grading exams, and hiring course readers and other teaching assistants.

ADVISING AND MENTORING

- **Anderson Truong** Sep 2025 – Present
Stanford University M.S. expected 2027
 - Anderson is investigating next-generation performance modeling of distributed, heterogeneous hardware.
- **Junwha Hong** May 2025 – Present
Agency for Defense Development South Korea
 - Junwha is researching ideas in scheduling agentic AI workloads on disaggregated hardware in collaboration with Professor Jongse Park.
- **Sai Gautham Ravipati** Jan 2024 – Present
Stanford University M.S. 2025, now Stanford Ph.D. expected 2030
 - Gautham is looking into compilation techniques for automatically generating tilings across memory hierarchies that target host-accelerator systems. We have already deployed this compilation flow to automatically generate tilings for testing the Onyx CGRA in [4].
- **Gina Sohn** Sep 2024 – Present
Stanford University Ph.D. expected 2027
 - Gina is developing dataflow abstractions and their compilers for dynamic tensor workloads, which are tensor computations where the tensor dimensions may be unknown or data-dependent. Example real-world applications include mixture-of-experts and dynamic runtime scheduling in deep learning models.
- **Genghan Zhang** Feb 2022 – Present
Tsinghua University B.S. 2023, now Stanford Ph.D. expected 2029
 - Genghan designed a compiler for optimized sparse tensor operations that detects the need for and inserts sparse temporary tensors (intermediate workspaces). This work achieved up to a 27x speedup and improved memory efficiency over dense temporary tensors while maintaining adaptability across various tensor expressions and user-defined implementations. This undergraduate research led to a publication at PLDI 2024 [9].
- **Rubens Lacouture** Sep 2021 – Present
Stanford University Ph.D. expected 2026
 - Rubens is currently looking into the compilation of sparse machine learning applications to dataflow backends, particularly focusing on extending ideas from the sparse abstract machine [13] to encompass machine learning operations [25] and the automatic fusion of distinct expressions to generated fused dataflow graphs [1].
- **Varun Desai** Jan 2024 – June 2025
Stanford University M.S. 2025

- Varun is currently looking into strategies for estimating the number of operations of a sparse workload given unknown input data.
- Varun also helped on the GPU evaluation of the Stardust compiler [7].
- **Manya Bansal** Feb 2022 – July 2023
Stanford University B.S. 2023, now MIT Ph.D. expected 2029 University
 - Manya developed Mosaic, an extensible compiler enabling flexible integration of external, high-performance libraries for tensor algebra, achieving up to a $206\times$ speedup by combining custom calls with efficient fused code generation. This undergraduate research led to a publication at PLDI 2023 and a distinguished paper award [12].
- **Zhouhua Xie** Jun 2023 – Present
Stanford University B.S. and M.S. expected 2025
 - Zhouhua is now looking into an FPGA emulation platform for the Onyx CGRA [10, 8].
 - Zhouhua investigated the performance bottlenecks and bugs in the sparse applications running on the Onyx CGRA [8, 10] and updated RTL implementations of the sparse primitives to fix those performance bottlenecks.
- **Parthiv Krishna** Sep 2022 – May 2024
Stanford University B.S. 2023, Stanford University M.S. expected 2024
 - Parthiv looked into implementing the sparse abstract machine primitives from [10, 8] on an FPGA.
- **Jayashree Adivarahan** Jun 2023 – Sep 2024
Arizona State University B.S. expected 2025
 - Jayashree helped with the evaluation of the Onyx CGRA [10, 8], specifically developing scripts that process and tile input tensor data before sending that data to the CGRA chip.
 - Jayashree looked into implementing the sparse abstract machine primitives from [10, 8] on an FPGA.
- **Akhilesh Balasingam** Jun 2023 – Feb 2024
Stanford University B.S. expected 2025
 - Akhilesh helped with the evaluation of the Onyx CGRA [10, 8], specifically running timing measurements on the physical CGRA and developing scripts that process and tile input tensor data.
- **Gloria Tumushabe** Sep 2016 – Mar 2023
UC Berkeley B.S. 2020, UC Berkeley M.S. 2021, Now Salesforce Senior Software Engineer
 - Gloria implemented techniques that help students form optional study groups in large-scale EECS courses at UC Berkeley that take into account inclusivity, which led to the Master's thesis "Inclusive and Scalable Study Group Formation" and a publication at SIGCSE TS 2023 [14].

HONORS AND AWARDS

-
- | | |
|---|---|
| • Rising Stars in EECS Awardee
MIT | 2024
[🔗] |
| • NextProf Nexus Future Faculty Awardee
University of Michigan, University of California Berkeley, and Georgia Tech | 2024 |
| • Distinguished Paper Award for Mosaic [12]
PLDI | 2023
[🔗] |
| • Alton B. Zerby and Carol T. Koerner Outstanding Student Award
IEEE–Eta Kappa Nu (HKN) | 2019
[📺] [🔗] [🔗] |
| • Graduate Research Fellowship (GRFP)
National Science Foundation (NSF) | 2019 |
| • Arthur M. Hopkin Award
University of California, Berkeley EECS Department | 2018–2019
[🔗] |
| • Outstanding Graduate Student Instructor (OGSI) Award
University of California, Berkeley | 2017–2018
[🔗] |
| • National Electrical Engineering Honor Society Member and Officer
IEEE–Eta Kappa Nu (HKN) | 2016 |
| • National Engineering Honor Society Member and Officer
Tau Beta Pi (TBP) | 2016 |
| • Leadership Award
Cal Alumni Association | 2015–2016
[🔗] |

WORK EXPERIENCE

- **Numenta** Jan 2025 – Mar 2026
Hardware Architecture Consultant Menlo Park, USA
- **Samsung Semiconductor** Jun 2023 – Sep 2023
Hardware Architecture Intern San Jose, USA
 - Researching the programmability and characterization of sparse workloads on next-generation supercomputing systems.
- **Ayar Labs** Jun 2020 – Sep 2020
Hardware VLSI Intern Santa Clara, USA
 - In charge of updating the design of multiple digital circuit blocks in System Verilog to match provided specifications, digital verification of those blocks, and pushing those blocks through the physical design process until they were met timing and were DRC and LVS clean.
- **Apple** Jun 2019 – Sep 2019, May 2018 – Aug 2018
CPU Design Verification Intern Cupertino, USA
 - Created instruction set architecture pseudocode for the Apple CPU in 2019.
 - Developed a runtime co-simulation checker using and added debug features into a C++ functional model of the Apple CPU in 2018.
- **NASA Jet Propulsion Laboratory (JPL)** Jun 2016 – Aug 2016
Hardware Research Intern Pasadena, USA
 - Developed a sensor testbed that verified the communication systems on a radiation-hardened, deep-space CubeSat avionics board.

TALKS

- **From Language to Silicon: Programming Systems for Sparse Accelerators** Dec 2025
Apple Inc
- **From Language to Silicon: Programming Systems for Sparse Accelerators** Dec 2025
Apple Inc
- **Portable Compilation of the Collection IR** Nov 2025
PRISM SRC Affiliates Meeting
- **From Language to Silicon: Programming Systems for Sparse Accelerators** Aug 2025
Stanford Portal Center Research Retreat
- **From Language to Silicon: Programming Systems for Sparse Accelerators** Feb – Apr 2025
Columbia University
University of Texas, Austin (UT Austin)
University of Michigan, Ann Arbor
Princeton University
Cornell University and Cornell Tech
University of Washington
University of California, Berkeley (UCB)
Massachusetts Institute of Technology (MIT)
University of Wisconsin, Madison (UW–Madison)
Carnegie Mellon University (CMU)
University of Maryland, College Park (UMD)
University of California, Santa Barbara (UCSB)
- **Mapping Sparse Applications to Accelerated Computing Systems** [📺]
Jun 2024
PLDI 2024: Sparse Workshop
- **The Sparse Abstract Machine** Apr 2024
Stanford and U-Tokyo Workshop: Agile and low-cost HW design workshop for domain-specific computing era
- **Architectural and compiler support for accelerating embedding operations on general purpose processors** Feb 2024
MIT Computer Science and Artificial Intelligence Laboratory (CSAIL)
- **Beyond the Sparse Abstract Machine: Mapping sparse applications to accelerated computing systems** Dec 2023
UC Santa Barbara Computer Science Department Colloquium
- **The Sparse Abstract Machine** Jun 2023
PLDI 2023: Compilation Techniques for Sparse Tensor Algebra (CTSTA)
- **Sparse Abstract Machine and Dataflow Hardware** Nov 2022
NVIDIA Tutorial on Sparse Tensor Algebra Compilation
- **The Sparse Abstract Machine: Sparse Tensor Algebra as Dataflow Graphs** Oct 2022
University of Washington: SAMPL Lunch Talk
- **The Sparse Abstract Machine: Sparse Tensor Algebra as Dataflow Graphs** Sep 2022
Samsung Semiconductor System Architecture Lab (SAL)
- **The Sparse Abstract Machine: Sparse Tensor Algebra as Dataflow Graphs** Aug 2022
AHA Agile Hardware Project Retreat
- **Compilation of Sparse Array Programming Models** Apr 2022
Google Brain: ML compilation seminar
- **Compilation of Sparse Array Programming Models** Apr 2022
University of California, Santa Cruz: LSD Seminar

PROFESSIONAL ACTIVITIES

- **ASPLOS 2027** 2026
ERC Member
- **TACO 2026** 2026
Reviewer
- **YArch Workshop at ISCA 2026** 2026
Co-Organizer and TPC Member
- **DAC 2026** 2025 – 2026
TPC Member
- **ASPLOS 2026** 2025
TPC Member (Summer)
- **CGO 2026** 2025
Artifact Evaluation Co-Chair and TPC Member
- **SPICE Workshop at MICRO 2025** 2025
Co-Organizer and TPC Member
- **LATTE Workshop at ASPLOS 2024** 2024
TPC Member
- **SIGGRAPH 2024** 2024
External Reviewer
- **Engineering Students for Diversity, Equity, and Inclusion (ES4DEI)** 2023 – 2024
College Track Co-Lead
- **Stanford University CS Department** 2023
PhD Admissions Committee
- **Stanford CS Department Student-Applicant Support Program** 2020 – 2021
Volunteer Reviewer
- **AfroFemCoders** 2020 – 2021
Advisor, Co-founder, and Instructor
- **UC Berkeley EECS Department Inaugural Women's History Month** 2019
Originator and Organizer

REFERENCES

Kunle Olukotun

Professor in Electrical Engineering and Computer Science
Stanford University
kunle@stanford.edu
Relationship: Primary Thesis Advisor

Joel Emer

Professor of the Practice in Electrical Engineering
and Computer Science, Researcher
MIT, NVIDIA
emer@csail.mit.edu
Relationship: Thesis Committee Member and Research Mentor

Fredrik Kjolstad

Assistant Professor in Computer Science
Stanford University
kjolstad@stanford.edu
Relationship: Thesis Co-advisor

Mark Horowitz

Professor in Electrical Engineering and Computer Science
Stanford University
horowitz@ee.stanford.edu
Relationship: Research Mentor